


Held out from training, across p : only the agents that could learn a proxy lose optimality


5×5 — proxy is free

trained $p=1.0$




11×11 — proxy costs a detour

trained $p=1.0$



11×11 control — no proxy to learn

trained $p=0.5$



p at test time (P that colour marks the higher-value objective)