


Held out from training, across ρ : only the agents that could learn a proxy lose optimality


5×5 — proxy is free

trained $\rho=1.0$




11×11 — proxy costs a detour

trained $\rho=1.0$



11×11 control — no proxy to learn

trained at chance, $\rho=0.5$



ρ at test time (P that colour marks the higher-value objective)