

Route transformer trained by next-token prediction: one line per seed

trained at $p=1.0$

trained at $p=0.5$ (control)

